

In [1]: ▶

```
1 import pandas as pd
2 import statsmodels as sm
3 import matplotlib.pyplot as plt
4 import statsmodels.formula.api as smf
5 import numpy as np
6 import seaborn as sns
7 import scipy.stats as stats
8 import yfinance as yf
9 from yfinance import download
10 import datetime
11 from datetime import datetime as dt
12 import warnings
13 warnings.filterwarnings('ignore')
14 from statsmodels.tsa.seasonal import seasonal_decompose
15 plt.style.use('seaborn-white')
16 %matplotlib inline
17 from statsmodels.tsa.stattools import adfuller
18 from statsmodels.graphics.tsaplots import plot_acf
19 from statsmodels.graphics.tsaplots import plot_pacf
20 from scipy.stats import gaussian_kde
21 from statsmodels.tsa.api import VAR
22 from statsmodels.tsa.holtwinters import ExponentialSmoothing as HWES
23 from statsmodels.tsa.seasonal import STL
24 from statsmodels.tsa.seasonal import seasonal_decompose
```

Problem 1

In [2]: ▶

1

housing_LA = pd.read_csv('LXXRNSA.csv', parse_dates = True, index_col

2

housing_LA = housing_LA[:-30] *#We already set aside the last 30 observ*

3

housing_LA

Out[2]:

LXXRNSA	
DATE	
1987-01-01	59.330841
1987-02-01	59.645596
1987-03-01	59.986172
1987-04-01	60.805706
1987-05-01	61.670846
...	...
2020-04-01	295.732705
2020-05-01	296.480456
2020-06-01	297.740428
2020-07-01	301.109304
2020-08-01	305.292342

404 rows × 1 columns

```
In [3]: 1 housing_SF = pd.read_csv('SFXRSA.csv', parse_dates = True, index_col =
2 housing_SF = housing_SF[:-30] #We already set aside the last 30 observ
3 housing_SF
```


Out[3]:

SFXRSA

DATE	
1987-01-01	46.955792
1987-02-01	47.302675
1987-03-01	47.840213
1987-04-01	47.984058
1987-05-01	48.305065
...	...
2020-04-01	271.616585
2020-05-01	270.924749
2020-06-01	270.802327
2020-07-01	273.929576
2020-08-01	278.827769

404 rows × 1 columns

```
In [4]: 1 fit,ax = plt.subplots(figsize = (15,8))
2 ax.plot(housing_LA, label = "Housing Price Index - Los Angeles")
3 ax.plot(housing_SF, label = "Housing Price Index - San Francisco")
4 ax.set_title("Housing Price Indices: Los Angeles vs San Francisco")
5 ax.set_ylabel("Index")
6 ax.set_xlabel("Year")
7 ax.legend()
8 ax.grid()
```



```
In [5]: 1 adfuller(housing_LA, regression='ct')
```

```
Out[5]: (-2.5321699589071507,  
0.3120694758121877,  
17,  
386,  
{'1%': -3.98241675663651,  
'5%': -3.421925528129767,  
'10%': -3.1337751777180234},  
869.101375915571)
```


```
In [6]: 1 adfuller(housing_SF, regression='ct')
```

```
Out[6]: (-2.3381658893083794,  
0.41296526399871225,  
4,  
399,  
{'1%': -3.9816401523738554,  
'5%': -3.4215509815220897,  
'10%': -3.1335552071923267},  
946.8796796873935)
```

based on AD test, it can be said that both data set are not stationary


```
In [7]: 1 plot_acf(housing_LA).suptitle('')
```

```
Out[7]: Text(0.5, 0.98, '')
```




```
In [8]: 1 plot_pacf(housing_LA).suptitle('')
```

```
Out[8]: Text(0.5, 0.98, '')
```




```
In [9]: 1 plot_acf(housing_SF).suptitle('')
```

```
Out[9]: Text(0.5, 0.98, '')
```



```
In [10]: 1 plot_pacf(housing_SF).suptitle('')
```

```
Out[10]: Text(0.5, 0.98, '')
```



```
In [11]: 1 #can change to percentage change
2 #just make sure it is stationary
3 housing_LA['diff_LA'] = housing_LA['LXXRNSA'].diff()
4 housing_LA.dropna(inplace = True)
5 housing_SF['diff_SF'] = housing_SF['SFXRSA'].diff()
6 housing_SF.dropna(inplace = True)
```

```
In [12]: 1 housing_LA['diff_LA'] = housing_LA['LXXRNSA'].diff()
2 housing_LA.dropna(inplace = True)
3 housing_SF['diff_SF'] = housing_SF['SFXRSA'].diff()
4 housing_SF.dropna(inplace = True)
```

```
In [13]: 1 #check statinarity again
```

```
In [14]: 1 adfuller(housing_LA['diff_LA'])
```

```
Out[14]: (-3.334884219140001,
0.013389360921523109,
16,
385,
{'1%': -3.4474498334928687,
'5%': -2.8690765390453703,
'10%': -2.570784795075055},
871.5455873313315)
```

In [15]: 1 adfuller(housing_SF['diff_SF'])

Out[15]: (-4.194205071337011,
0.0006739417628468098,
11,
390,
{'1%': -3.4472291365835566,
'5%': -2.8689795375849223,
'10%': -2.5707330834976987},
946.9195064178084)

In [16]: 1 *#estimate the model ; 13 lags is optimal according to model*

```
In [17]: ▶ 1 VAR_variables = pd.concat([housing_LA['diff_LA'], housing_SF['diff_SF']  
2  
3 # Fit the VAR model  
4 model_VAR = VAR(VAR_variables)  
5 results = model_VAR.fit(maxlags=15, ic='aic')  
6  
7 # Print the summary of the VAR model results  
8 print(results.summary())
```


Summary of Regression Results

```
=====
Model:                VAR
Method:               OLS
Date:                Thu, 25, May, 2023
Time:                22:20:23
```

```
-----
No. of Equations:      2.00000      BIC:                -0.556161
Nobs:                 388.000      HQIC:               -0.913508
Log likelihood:       -820.332     FPE:                0.317362
AIC:                 -1.14827     Det(Omega_mle):     0.274756
-----
```

Results for equation diff_LA

```
=====
=====
```

	coefficient	std. error	t-stat
prob			

const	0.039754	0.039318	1.011
0.312			
L1.diff_LA	0.895635	0.054440	16.452
0.000			
L1.diff_SF	0.184084	0.050113	3.673
0.000			
L2.diff_LA	0.022486	0.072833	0.309
0.758			
L2.diff_SF	-0.086496	0.062101	-1.393
0.164			
L3.diff_LA	-0.152500	0.072275	-2.110
0.035			
L3.diff_SF	0.045209	0.063025	0.717
0.473			
L4.diff_LA	0.203394	0.074469	2.731
0.006			
L4.diff_SF	-0.032300	0.067451	-0.479
0.632			
L5.diff_LA	-0.135936	0.075824	-1.793
0.073			
L5.diff_SF	-0.017634	0.069245	-0.255
0.799			
L6.diff_LA	-0.057693	0.075945	-0.760
0.447			
L6.diff_SF	0.042414	0.069139	0.613
0.540			
L7.diff_LA	-0.016848	0.076108	-0.221
0.825			
L7.diff_SF	0.103080	0.069599	1.481
0.139			
L8.diff_LA	0.005864	0.075975	0.077
0.938			
L8.diff_SF	-0.181096	0.070077	-2.584
0.010			
L9.diff_LA	0.018359	0.076148	0.241
0.809			
L9.diff_SF	0.108314	0.070124	1.545
0.122			

L10.diff_LA 0.217	0.093263	0.075565	1.234
L10.diff_SF 0.286	-0.075311	0.070516	-1.068
L11.diff_LA 0.000	0.275387	0.074546	3.694
L11.diff_SF 0.721	-0.024757	0.069419	-0.357
L12.diff_LA 0.050	-0.145142	0.073930	-1.963
L12.diff_SF 0.036	-0.136749	0.065228	-2.096
L13.diff_LA 0.234	-0.087268	0.073399	-1.189
L13.diff_SF 0.376	0.057827	0.065268	0.886
L14.diff_LA 0.521	-0.035527	0.055409	-0.641
L14.diff_SF 0.134	0.080100	0.053416	1.500

=====

=====

Results for equation diff_SF

=====

=====

	coefficient	std. error	t-stat
prob			

const 0.066	0.078554	0.042769	1.837
L1.diff_LA 0.131	0.089451	0.059219	1.511
L1.diff_SF 0.000	0.789341	0.054512	14.480
L2.diff_LA 0.463	-0.058135	0.079226	-0.734
L2.diff_SF 0.002	0.208064	0.067552	3.080
L3.diff_LA 0.001	0.272026	0.078619	3.460
L3.diff_SF 0.000	-0.490323	0.068557	-7.152
L4.diff_LA 0.061	-0.151710	0.081006	-1.873
L4.diff_SF 0.001	0.254544	0.073372	3.469
L5.diff_LA 0.624	-0.040462	0.082480	-0.491
L5.diff_SF 0.636	0.035604	0.075323	0.473
L6.diff_LA 0.113	0.130821	0.082611	1.584
L6.diff_SF 0.065	-0.138726	0.075207	-1.845
L7.diff_LA	-0.057760	0.082789	-0.698

0.485			
L7.diff_SF	0.065021	0.075708	0.859
0.390			
L8.diff_LA	0.015896	0.082643	0.192
0.847			
L8.diff_SF	0.003699	0.076228	0.049
0.961			
L9.diff_LA	0.014738	0.082832	0.178
0.859			
L9.diff_SF	0.055511	0.076280	0.728
0.467			
L10.diff_LA	-0.006676	0.082198	-0.081
0.935			
L10.diff_SF	-0.068960	0.076705	-0.899
0.369			
L11.diff_LA	0.213491	0.081090	2.633
0.008			
L11.diff_SF	0.002643	0.075512	0.035
0.972			
L12.diff_LA	-0.079015	0.080419	-0.983
0.326			
L12.diff_SF	-0.174254	0.070953	-2.456
0.014			
L13.diff_LA	-0.307906	0.079842	-3.856
0.000			
L13.diff_SF	0.125995	0.070997	1.775
0.076			
L14.diff_LA	0.185274	0.060273	3.074
0.002			
L14.diff_SF	-0.016674	0.058104	-0.287
0.774			

=====

=====

Correlation matrix of residuals

	diff_LA	diff_SF
diff_LA	1.000000	0.264553
diff_SF	0.264553	1.000000

Problem 2

```
In [18]: ▶ 1 # granger causality, look at p-value
          2 # we are testing if LA has effect on SF. Low p-value signifies La DOES
```

```
In [19]: 1 results.test_causality('diff_LA','diff_SF', kind='f').summary()
```

Out[19]: Granger causality F-test. H_0: diff_SF does not
Granger-cause diff_LA. Conclusion: reject H_0 at 5%
significance level.

Test statistic	Critical value	p-value	df
3.994	1.706	0.000	(14, 718)

The test statistic is 3.994, and the critical value is 1.706. The p-value is 0.000, which is less than the significance level of 0.05. Therefore, we have sufficient evidence to conclude that there is Granger causality between the two series diff_SF and diff_LA

```
In [20]: 1 #then we test the other way around
          2 #Low
```

```
In [21]: 1 results.test_causality('diff_SF', 'diff_LA', kind='f').summary()
```

Out[21]: Granger causality F-test. H_0: diff_LA does not
Granger-cause diff_SF. Conclusion: reject H_0 at 5%
significance level.

Test statistic	Critical value	p-value	df
4.540	1.706	0.000	(14, 718)

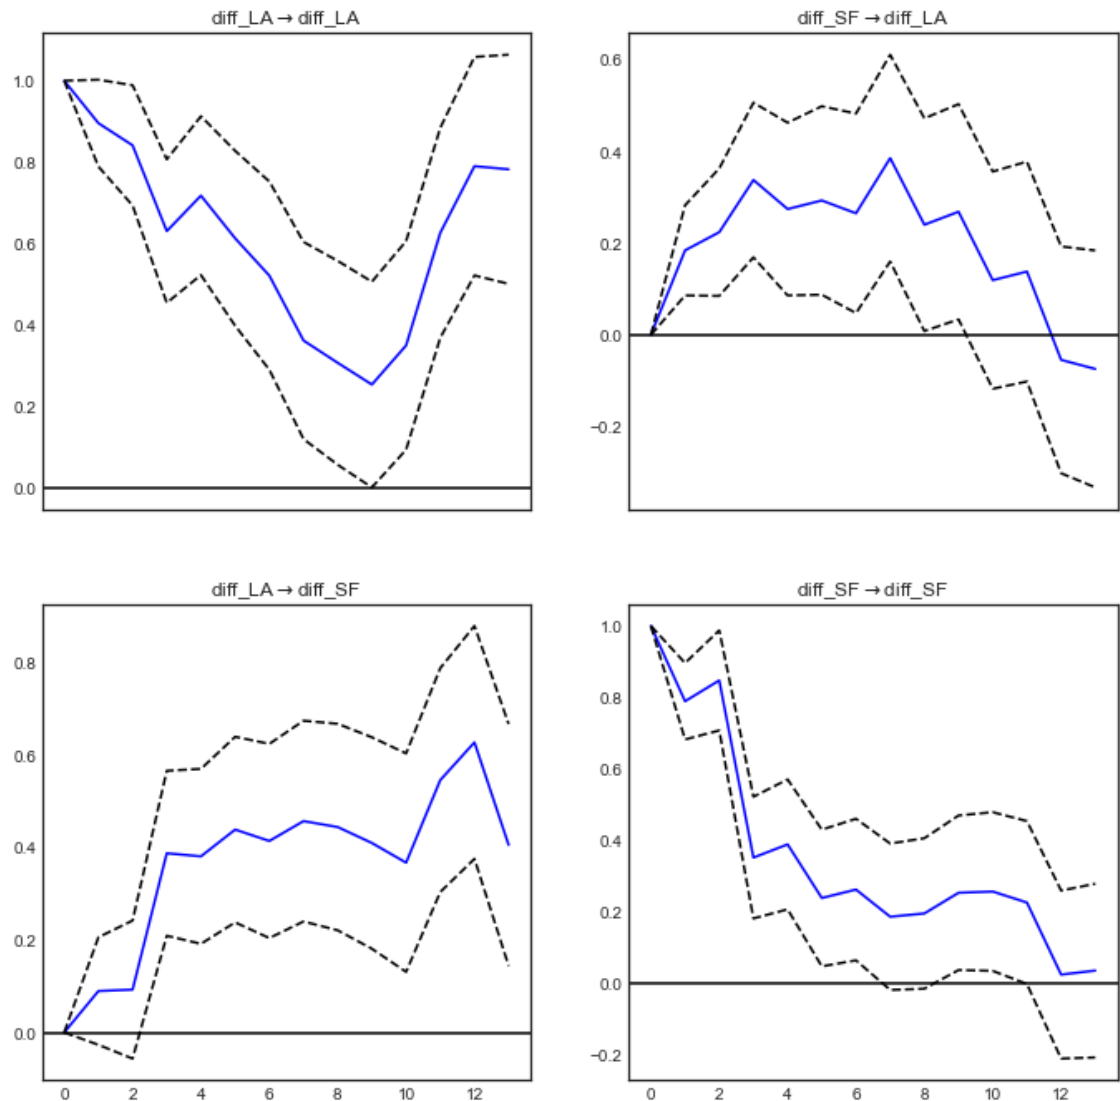
The test statistic is 4.540, and the critical value is 1.706. The p-value is 0.000, which is less than the significance level of 0.05. Therefore, we have sufficient evidence to conclude that there is Granger causality between the two series diff_LA and diff_SF

Problem 3

```
In [22]: 1 #3 when band reaches 0 we can say it is not SS; should not cross the 0
          2 #shock in SF effect on LA. if you shock prices in SF is has effect in
```

```
In [23]: 1 # IRFs
2 irf = results.irf(13)
3 irf.plot(orth=False).suptitle('') # Plots all of them
```

Out[23]: Text(0.5, 0.98, '')



Based on the above output, it can be observed that there is statistically significant Granger causality between the context of housing price index changes in Los Angeles and San Francisco.

Specifically, the top-right panel suggests that shocks in the change in the price index of housing in Los Angeles have a statistically significant impact on the change in the price index of housing in San Francisco, lasting up to 4 lags. Similarly, the bottom-left panel indicates that shocks in the change in the price index of housing in San Francisco have a positive impact on itself, lasting up to 5 lags.

Overall, these findings imply that there is a relationship between the housing markets of Los Angeles and San Francisco, where changes in one market can influence the other.

Problem 4

In [8]:

▶

1

sales only

In [69]:

▶

1

import pandas as pd

2

3

retaildata = pd.read_excel("retail.xlsx", skiprows=1)

4

retaildata

Out[69]:

	Series ID	A3349335T	A3349627V	A3349338X	A3349398A	A3349468W	A3349336V	A334
0	1982-04-01	303.1	41.7	63.9	408.7	65.8	91.8	
1	1982-05-01	297.8	43.1	64.0	404.9	65.8	102.6	
2	1982-06-01	298.0	40.3	62.7	401.0	62.3	105.0	
3	1982-07-01	307.9	40.9	65.6	414.4	68.2	106.0	
4	1982-08-01	299.2	42.1	62.6	403.8	66.0	96.9	
...	
376	2013-08-01	2244.2	264.6	247.8	2756.6	305.0	423.2	
377	2013-09-01	2157.0	262.8	240.2	2660.1	292.1	401.3	
378	2013-10-01	2299.5	264.4	244.5	2808.4	342.3	401.1	
379	2013-11-01	2271.3	271.5	232.2	2775.1	359.0	444.0	
380	2013-12-01	2612.8	394.5	270.9	3278.2	427.0	667.2	

381 rows × 190 columns

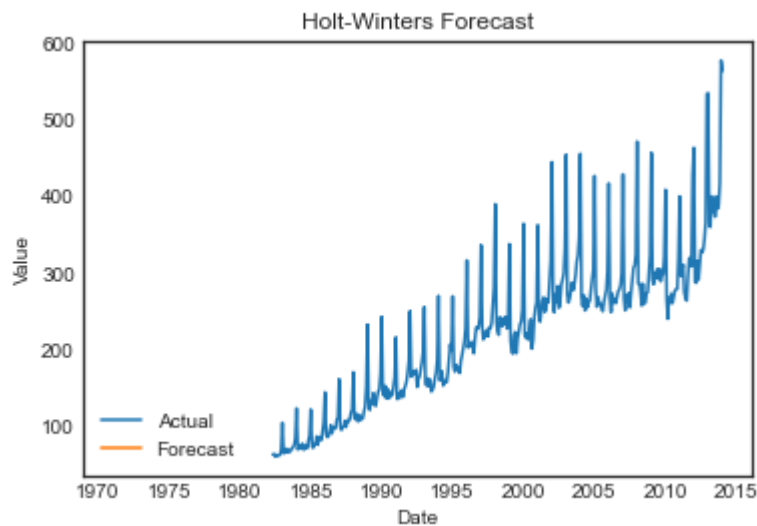

◀

▶

```
In [70]: ▶ 1 # import pandas as pd
          2 # import numpy as np
          3 # import matplotlib.pyplot as plt
          4
          5 # # Assuming "your_column_name" is the column name you want to select
          6 # selected_column = retaildata["A3349873A"]
          7
          8 # # Create a time series with a frequency of 12 (monthly) starting from
          9 # myts = pd.Series(selected_column.values, pd.date_range(start="1982-01-01",
10
11 # # Plot the time series
12 # plt.plot(myts)
13 # plt.xlabel("Date")
14 # plt.ylabel("Value")
15 # plt.title("Time Series")
16 # plt.show()
17
```

```
In [71]: ▶ 1 # retaildata_train = retaildata[:-1]
          2 # retaildata_test = retaildata[-1:]
```

```
In [72]: 1 import pandas as pd
2 import matplotlib.pyplot as plt
3 from statsmodels.tsa.holtwinters import ExponentialSmoothing
4
5 # Assuming "your_column_name" is the column name you want to select
6 selected_column = retaildata["A3349873A"]
7
8 # Create a time series with a frequency of 12 (monthly) starting from
9 myts = pd.Series(selected_column.values, pd.date_range(start="1982-04",
10
11 # Plot the time series
12 plt.plot(myts)
13 plt.xlabel("Date")
14 plt.ylabel("Value")
15 plt.title("Time Series")
16 plt.show()
17
18 retaildata_train = retaildata[:-1]
19 retaildata_test = retaildata[-1:]
20
21 # Select the specific column as a Series
22 retaildata_train_series = retaildata_train['A3349873A']
23 retaildata_test_series = retaildata_test['A3349873A']
24
25 # Convert the data to a numeric type
26 retaildata_train_series = pd.to_numeric(retaildata_train_series, errors='coerce')
27 retaildata_test_series = pd.to_numeric(retaildata_test_series, errors='coerce')
28
29 # Fit the Holt-Winters model to the training data
30 hw_model = ExponentialSmoothing(retaildata_train_series, trend='mul',
31 hw_fitted = hw_model.fit()
32
33 # Generate the forecasts
34 forecast = hw_fitted.predict(start=retaildata_test_series.index[0], end=retaildata_test_series.index[-1])
35
36 # Visualize the results
37 plt.plot(myts, label='Actual')
38 plt.plot(forecast, label='Forecast')
39 plt.xlabel("Date")
40 plt.ylabel("Value")
41 plt.title("Holt-Winters Forecast")
42 plt.legend()
43 plt.show()
44
```

```
In [4]: 1 # hw_model = HWES(retaildata_train, seasonal_periods=12, trend = 'mul')
        2 # hw_fitted = hw_model.fit()
```

In [74]: `hw_fitted.summary()`


Out[74]: ExponentialSmoothing Model Results

Dep. Variable:	A3349873A	No. Observations:	380
Model:	ExponentialSmoothing	SSE	66754.210
Optimized:	True	AIC	1996.069
Trend:	Multiplicative	BIC	2059.111
Seasonal:	Multiplicative	AICC	1997.963
Seasonal Periods:	12	Date:	Thu, 25 May 2023
Box-Cox:	False	Time:	22:20:28
Box-Cox Coeff.:	None		

	coeff	code	optimized
smoothing_level	0.5546165	alpha	True
smoothing_trend	5.4668e-11	beta	True
smoothing_seasonal	0.4453835	gamma	True
initial_level	50.191696	l.0	True
initial_trend	1.0029363	b.0	True
initial_seasons.0	1.2395683	s.0	True
initial_seasons.1	1.2923666	s.1	True
initial_seasons.2	1.2079926	s.2	True
initial_seasons.3	1.2731564	s.3	True
initial_seasons.4	1.3064631	s.4	True
initial_seasons.5	1.3366330	s.5	True
initial_seasons.6	1.4054417	s.6	True
initial_seasons.7	1.4732702	s.7	True
initial_seasons.8	2.1516625	s.8	True
initial_seasons.9	1.1986225	s.9	True
initial_seasons.10	1.1739294	s.10	True
initial_seasons.11	1.2595666	s.11	True

```
In [75]: 1 hw_fitted.fittedvalues.plot()
```

Out[75]: <AxesSubplot:>



```
In [76]: 1 #you have to compare rmse with 2 methods; thats why you need 2 models
```

```
In [77]: 1 # Convert the data to a numeric type
2 retaildata_train_series = pd.to_numeric(retaildata_train_series, error
3
4 # Fit the Holt-Winters model to the training data with damped trend
5 hw_model_damped_trend = ExponentialSmoothing(retaildata_train_series,
6 hw_fitted_damped_trend = hw_model_damped_trend.fit(damping_trend=0.5)
```

In [78]: `1 hw_fitted_damped_trend.summary()`


Out[78]: ExponentialSmoothing Model Results

Dep. Variable:	A3349873A	No. Observations:	380
Model:	ExponentialSmoothing	SSE	67265.782
Optimized:	True	AIC	2000.970
Trend:	Multiplicative	BIC	2067.952
Seasonal:	Multiplicative	AICC	2003.081
Seasonal Periods:	12	Date:	Thu, 25 May 2023
Box-Cox:	False	Time:	22:20:28
Box-Cox Coeff.:	None		

	coeff	code	optimized
smoothing_level	0.5586861	alpha	True
smoothing_trend	1.2992e-11	beta	True
smoothing_seasonal	0.4413139	gamma	True
initial_level	67.319267	l.0	True
initial_trend	0.9352466	b.0	True
damping_trend	0.5000000	phi	False
initial_seasons.0	0.9911108	s.0	True
initial_seasons.1	1.0413237	s.1	True
initial_seasons.2	0.9720482	s.2	True
initial_seasons.3	1.0281134	s.3	True
initial_seasons.4	1.0567893	s.4	True
initial_seasons.5	1.0832042	s.5	True
initial_seasons.6	1.1422145	s.6	True
initial_seasons.7	1.2001018	s.7	True
initial_seasons.8	1.7583680	s.8	True
initial_seasons.9	0.9739911	s.9	True
initial_seasons.10	0.9559678	s.10	True
initial_seasons.11	1.0273874	s.11	True

```
In [79]: 1 hw_fitted_damped_trend.fittedvalues.plot()
```

Out[79]: <AxesSubplot:>



```
In [80]: 1 hw_fitted_forecast = hw_fitted.forecast(1)
```

```
In [81]: 1 hw_model_damped_trend = HWES(retaildata_train_series, seasonal_periods=
2 hw_fitted_damped_trend = hw_model_damped_trend.fit(damping_trend = 0.5
3 hw_fitted_damped_trend.summary()
```

Out[81]:


ExponentialSmoothing Model Results

Dep. Variable:	A3349873A	No. Observations:	380
Model:	ExponentialSmoothing	SSE	67265.782
Optimized:	True	AIC	2000.970
Trend:	Multiplicative	BIC	2067.952
Seasonal:	Multiplicative	AICC	2003.081
Seasonal Periods:	12	Date:	Thu, 25 May 2023
Box-Cox:	False	Time:	22:20:29
Box-Cox Coeff.:	None		

	coeff	code	optimized
smoothing_level	0.5586861	alpha	True
smoothing_trend	1.2992e-11	beta	True
smoothing_seasonal	0.4413139	gamma	True
initial_level	67.319267	l.0	True
initial_trend	0.9352466	b.0	True
damping_trend	0.5000000	phi	False
initial_seasons.0	0.9911108	s.0	True
initial_seasons.1	1.0413237	s.1	True
initial_seasons.2	0.9720482	s.2	True
initial_seasons.3	1.0281134	s.3	True
initial_seasons.4	1.0567893	s.4	True
initial_seasons.5	1.0832042	s.5	True
initial_seasons.6	1.1422145	s.6	True
initial_seasons.7	1.2001018	s.7	True
initial_seasons.8	1.7583680	s.8	True
initial_seasons.9	0.9739911	s.9	True
initial_seasons.10	0.9559678	s.10	True
initial_seasons.11	1.0273874	s.11	True

```
In [82]: 1 hw_fitted_damped_trend.fittedvalues.plot()
```

Out[82]: <AxesSubplot:>



```
In [83]: 1 hw_fitted_forecast = hw_fitted.forecast(1)
```

```
In [84]: 1 hw_fitted_damped_trend_forecast = hw_fitted_damped_trend.forecast(1)
```

```
In [85]: 1 hw_fitted_damped_trend_forecast = hw_fitted_damped_trend.forecast(1)
```

```
In [86]: 1 # Accuracy metrics
2 def rmse(forecast, actual):
3     rmse = np.mean((forecast - actual)**2)**.5 # RMSE
4     return({'rmse':rmse})
```

```
In [87]: 1 #Lower RMSe is better; downtrend
```

```
In [88]: 1 rmse(hw_fitted_forecast, aus_retail_agg_test['Turnover'][0])
```


Out[88]: {'rmse': 64435.83706890834}

```
In [89]: 1 rmse(hw_fitted_damped_trend_forecast, aus_retail_agg_test['Turnover'][(
```

Out[89]: {'rmse': 64438.67812117148}

```
In [90]: 1 hw_fitted_damped_trend.resid.plot()
```

```
Out[90]: <AxesSubplot:>
```



```
In [91]: 1 retaildata_train_series_RMSE = aus_retail_agg[:'2010-12-01']  
2 retaildata_test_series_RMSE = aus_retail_agg['2010-12-01':]
```

```
In [92]: 1 #training set now different; you need to run holt winters model again  
2 # did 12 bc we will have probelm with seasonal n approach forecast bas
```


In [93]:

```

1 hw_model_damped_trend2 = HWES(retaildata_train_series_RMSE, seasonal_p
2 hw_model_damped_trend2 = hw_model_damped_trend2.fit(damping_trend = 0.
3 hw_model_damped_trend2.summary()

```

Out[93]:

ExponentialSmoothing Model Results

Dep. Variable:	Turnover	No. Observations:	345
Model:	ExponentialSmoothing	SSE	64395129.744
Optimized:	True	AIC	4221.266
Trend:	Multiplicative	BIC	4286.607
Seasonal:	Multiplicative	AICC	4223.605
Seasonal Periods:	12	Date:	Thu, 25 May 2023
Box-Cox:	False	Time:	22:20:29
Box-Cox Coeff.:	None		

	coeff	code	optimized
smoothing_level	0.4168593	alpha	True
smoothing_trend	0.4168593	beta	True
smoothing_seasonal	0.2157184	gamma	True
initial_level	6625.1990	l.0	True
initial_trend	0.9908539	b.0	True
damping_trend	0.5000000	phi	False
initial_seasons.0	0.9554139	s.0	True
initial_seasons.1	1.0095743	s.1	True
initial_seasons.2	0.9491036	s.2	True
initial_seasons.3	0.9650470	s.3	True
initial_seasons.4	0.9567505	s.4	True
initial_seasons.5	0.9379636	s.5	True
initial_seasons.6	0.9763789	s.6	True
initial_seasons.7	1.0348839	s.7	True
initial_seasons.8	1.3507723	s.8	True
initial_seasons.9	0.9550792	s.9	True
initial_seasons.10	0.8913703	s.10	True
initial_seasons.11	0.9717872	s.11	True

```
In [94]: 1 hw_fitted_forecast_test = hw_model_damped_trend2.forecast(12)
2 column_series = pd.Series(retaildata_test_series_RMSE['Turnover'])[0:12]
3 column_series = column_series.reset_index(drop=True)
4 hw_fitted_forecast_test = hw_fitted_forecast_test.reset_index(drop=True)
```

```
In [95]: 1 rmse(hw_fitted_forecast_test, column_series)
```

```
Out[95]: {'rmse': 4351.127593630717}
```

```
In [96]: 1 #we can compare with first 12 observatins from data set
```

```
In [97]: 1 # Generate seasonal naive forecasts
```

```
In [98]: 1 seasonal_naive_forecasts = retaildata_train_series_RMSE[-12:]
2 seasonal_naive_forecasts.reset_index(drop = True, inplace = True)
3 seasonal_naive_forecasts_series = pd.Series(seasonal_naive_forecasts['Turnover'])
4 seasonal_naive_forecasts_series
```

```
Out[98]: 0    37917.0
1    33565.5
2    37139.7
3    36214.8
4    37192.7
5    36692.5
6    38400.7
7    37927.1
8    37939.9
9    39188.5
10   40133.4
11   49799.7
Name: Turnover, dtype: float64
```

```
In [99]: 1 rmse(seasonal_naive_forecasts_series, column_series)
```

```
Out[99]: {'rmse': 4529.10961521872}
```

```
In [100]: 1 # same data, do box cox ets, and stl on time sereis
```

```
In [101]: 1 transformed_data, lambda_val = stats.boxcox(retaildata_train_series_RM
2 transformed_data = pd.Series(transformed_data)
3 transformed_data = pd.DataFrame(transformed_data, columns=retaildata_t
4 index = retaildata_train_series_RMSE.reset_index()
5 transformed_data = pd.merge(transformed_data, index['Month'], left_inc
6 transformed_data = transformed_data.set_index('Month')
7 transformed_data
```


Out[101]:

Turnover	
Month	
1982-04-01	21.713310
1982-05-01	21.838213
1982-06-01	21.662093
1982-07-01	21.851107
1982-08-01	21.663389
...	...
2010-08-01	32.434574
2010-09-01	32.436928
2010-10-01	32.663477
2010-11-01	32.831040
2010-12-01	34.382712

345 rows × 1 columns

```
In [ ]: 1 ## part
```

```
In [102]: 1 #stl
2 stl = STL(transformed_data)
3 results = stl.fit()
4 fig1 = results.plot()
```



```
In [103]: 1 stl_result = seasonal_decompose(transformed_data['Turnover'], model='n
2 seasonal = stl_result.seasonal
3 seasonally_adjusted_data = transformed_data['Turnover']/seasonal
4 seasonally_adjusted_data
```

```
Out[103]: Month
1982-04-01    21.937848
1982-05-01    21.878834
1982-06-01    21.890430
1982-07-01    21.966054
1982-08-01    21.785023
...
2010-08-01    32.616686
2010-09-01    32.663646
2010-10-01    32.548208
2010-11-01    32.541603
2010-12-01    32.383276
Length: 345, dtype: float64
```

```
In [104]: 1 seasonally_adjusted_data.index.freq='MS'
2
3 ets_model=sm.tsa.statespace.exponential_smoothing.ExponentialSmoothing
4                                     trend=True,
5                                     initialization_method='heuristic',
6                                     seasonal=12,
7                                     damped_trend=False).fit()
```

In [105]: ▶ 1 ets_model.summary()

Out[105]:

Exponential Smoothing Results

Dep. Variable:	y	No. Observations:	345
Model:	ETS(A, A, A)	Log Likelihood	219.407
Date:	Thu, 25 May 2023	AIC	-430.813
Time:	22:20:30	BIC	-415.439
Sample:	04-01-1982	HQIC	-424.691
	- 12-01-2010	Scale	0.016
Covariance Type:	opg		

	coef	std err	z	P> z	[0.025	0.975]
smoothing_level	0.2674	0.039	6.791	0.000	0.190	0.345
smoothing_trend	0.0044	0.003	1.443	0.149	-0.002	0.010
smoothing_seasonal	0.2234	0.034	6.547	0.000	0.156	0.290

initialization method: heuristic

level	21.9996
trend	0.0396
seasonal	-0.0312
seasonal.L1	0.0230
seasonal.L2	-0.1644
seasonal.L3	0.2330
seasonal.L4	0.1084
seasonal.L5	-0.0641
seasonal.L6	-0.0995
seasonal.L7	0.0310
seasonal.L8	-0.0472
seasonal.L9	-0.1221
seasonal.L10	0.1723
seasonal.L11	-0.0392

Ljung-Box (L1) (Q):	13.73	Jarque-Bera (JB):	3.95
Prob(Q):	0.00	Prob(JB):	0.14
Heteroskedasticity (H):	0.56	Skew:	-0.15
Prob(H) (two-sided):	0.00	Kurtosis:	3.43

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [106]: 1 #then divide by seasonal; seasonal adjustment, after adjustment, you c
2 #when you have ets model, you can do the 12 step ahead.
3 #we also have to transform the test dataset otherwise you will have ab
```

```
In [107]: 1 forecast_ets_model = ets_model.get_forecast(steps=12)
2 forecast_ets_model = forecast_ets_model.predicted_mean
3 forecast_ets_model = forecast_ets_model.reset_index(drop=True)
```

```
In [108]: 1 transformed_test = stats.boxcox(aus_retail_agg_test_RMSE['Turnover'])[0]
2 transformed_test = pd.Series(transformed_test)
3 transformed_test = pd.DataFrame(transformed_test, columns=aus_retail_a
4 index = aus_retail_agg_test_RMSE.reset_index()
5 transformed_test = pd.merge(transformed_test, index['Month'], left_in
6 transformed_test = transformed_test.set_index('Month')
7 transformed_test
```

Out[108]:

Turnover	
Month	
2010-12-01	34.382712
2011-01-01	32.497644
2011-02-01	31.805962
2011-03-01	32.466845
2011-04-01	32.395218
2011-05-01	32.430656
2011-06-01	32.363435
2011-07-01	32.592290
2011-08-01	32.622624
2011-09-01	32.644888
2011-10-01	32.883934
2011-11-01	33.070872

```
In [109]: 1 column_series2 = transformed_test['Turnover'].reset_index(drop=True)
2 rmse(forecast_ets_model, column_series2)
```

Out[109]: {'rmse': 0.5986192680089508}

Problem 5

In [189]:

▶

```
1 transformed_data, lambda_val = stats.boxcox(aus_retail_agg_train_RMSE[
2 transformed_data = pd.Series(transformed_data)
3 transformed_data = pd.DataFrame(transformed_data, columns=aus_retail_a
4 index = aus_retail_agg_train_RMSE.reset_index()
5 transformed_data = pd.merge(transformed_data, index['Month'], left_inc
6 transformed_data = transformed_data.set_index('Month')
7 transformed_data
```


Out[189]:

Turnover

Month	
1982-04-01	21.713310
1982-05-01	21.838213
1982-06-01	21.662093
1982-07-01	21.851107
1982-08-01	21.663389
...	...
2010-08-01	32.434574
2010-09-01	32.436928
2010-10-01	32.663477
2010-11-01	32.831040
2010-12-01	34.382712

345 rows × 1 columns


```
In [190]: 1 stl = STL(transformed_data)
          2
          3 results = stl.fit()
          4
          5 fig1 = results.plot()
```



```
In [191]: 1 stl_result = seasonal_decompose(transformed_data['Turnover'], model='n
          2 seasonal = stl_result.seasonal
          3 seasonally_adjusted_data = transformed_data['Turnover']/seasonal
          4 seasonally_adjusted_data
```

```
Out[191]: Month
1982-04-01    21.937848
1982-05-01    21.878834
1982-06-01    21.890430
1982-07-01    21.966054
1982-08-01    21.785023
...
2010-08-01    32.616686
2010-09-01    32.663646
2010-10-01    32.548208
2010-11-01    32.541603
2010-12-01    32.383276
Length: 345, dtype: float64
```

```
In [192]: 1 seasonally_adjusted_data.index.freq='MS'
          2
          3 ets_model=sm.tsa.statespace.exponential_smoothing.ExponentialSmoothing
          4                                     trend=True,
          5                                     initialization_method='he
          6                                     seasonal=12,
          7                                     damped_trend=False).fit()
```

In [193]: ▶ 1 ets_model.summary()

Out[193]:

Exponential Smoothing Results

Dep. Variable:	y	No. Observations:	345
Model:	ETS(A, A, A)	Log Likelihood	219.407
Date:	Fri, 26 May 2023	AIC	-430.813
Time:	15:23:33	BIC	-415.439
Sample:	04-01-1982	HQIC	-424.691
	- 12-01-2010	Scale	0.016
Covariance Type:	opg		

	coef	std err	z	P> z	[0.025	0.975]
smoothing_level	0.2674	0.039	6.791	0.000	0.190	0.345
smoothing_trend	0.0044	0.003	1.443	0.149	-0.002	0.010
smoothing_seasonal	0.2234	0.034	6.547	0.000	0.156	0.290

initialization method: heuristic

level	21.9996
trend	0.0396
seasonal	-0.0312
seasonal.L1	0.0230
seasonal.L2	-0.1644
seasonal.L3	0.2330
seasonal.L4	0.1084
seasonal.L5	-0.0641
seasonal.L6	-0.0995
seasonal.L7	0.0310
seasonal.L8	-0.0472
seasonal.L9	-0.1221
seasonal.L10	0.1723
seasonal.L11	-0.0392

Ljung-Box (L1) (Q):	13.73	Jarque-Bera (JB):	3.95
Prob(Q):	0.00	Prob(JB):	0.14
Heteroskedasticity (H):	0.56	Skew:	-0.15
Prob(H) (two-sided):	0.00	Kurtosis:	3.43

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [194]: 1 forecast_ets_model = ets_model.get_forecast(steps=97)
          2 forecast_ets_model = forecast_ets_model.predicted_mean
          3 forecast_ets_model = forecast_ets_model.reset_index(drop=True)
```

```
In [195]: 1 transformed_test = stats.boxcox(aus_retail_agg_test_RMSE['Turnover'],
          2 transformed_test = pd.Series(transformed_test)
          3 transformed_test = pd.DataFrame(transformed_test, columns=aus_retail_a
          4 index = aus_retail_agg_test_RMSE.reset_index()
          5 transformed_test = pd.merge(transformed_test, index['Month'], left_in
          6 transformed_test = transformed_test.set_index('Month')
          7 transformed_test
```

Out[195]:

	Turnover
Month	
2010-12-01	34.382712
2011-01-01	32.497644
2011-02-01	31.805962
2011-03-01	32.466845
2011-04-01	32.395218
...	...
2018-08-01	34.513901
2018-09-01	34.427226
2018-10-01	34.792395
2018-11-01	35.090448
2018-12-01	36.386514

97 rows × 1 columns

```
In [196]: 1 column_series2 = transformed_test['Turnover'].reset_index(drop=True)
          2 rmse(forecast_ets_model, column_series2)
```

Out[196]: {'rmse': 0.6693108961781675}

the ETS (Error, Trend, Seasonality) model with a Box-Cox transformation and seasonal adjustments based on STL decomposition outperforms the Holt-Winters' multiplicative method without a damped trend. The Root Mean Square Error (RMSE) for the ETS model with Box-Cox and STL adjustments is significantly lower compared to the other method, indicating better accuracy and forecasting performance.

Looking at the model details, the ETS model with Box-Cox and STL adjustments shows favorable parameter estimates. The smoothing level (0.2674) and smoothing seasonal (0.2234) coefficients are statistically significant, suggesting the presence of trend and seasonality components in the data. However, the smoothing trend coefficient (0.0044) is not statistically significant, indicating a relatively flat trend.

The method for the ETS model is heuristic, and the estimated levels, trends, and seasonal coefficients provide insights into the time series patterns. The Ljung-Box test result indicates the absence of autocorrelation at lag 1, and the Jarque-Bera test suggests the approximate normality of the model residuals. Furthermore, the Heteroskedasticity test reveals a moderate level of heteroskedasticity, while the skewness and kurtosis values means a slightly away from normality.

But all in all. the performance. as indicated by the lower RMSE value. The model captures the

Problem 6

- done in R


In []: 1 *#we will now have to do 4 models*

In [175]: 1 *##part a*


Time-Series [1:240] from 1985 to 2005: 75.7 75.4 83.1 82.9 77.3 ...

```
head(visitors) May Jun Jul Aug Sep Oct 1985 75.7 75.4 83.1 82.9 77.3 105.7
```

```
In [177]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
In [178]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```

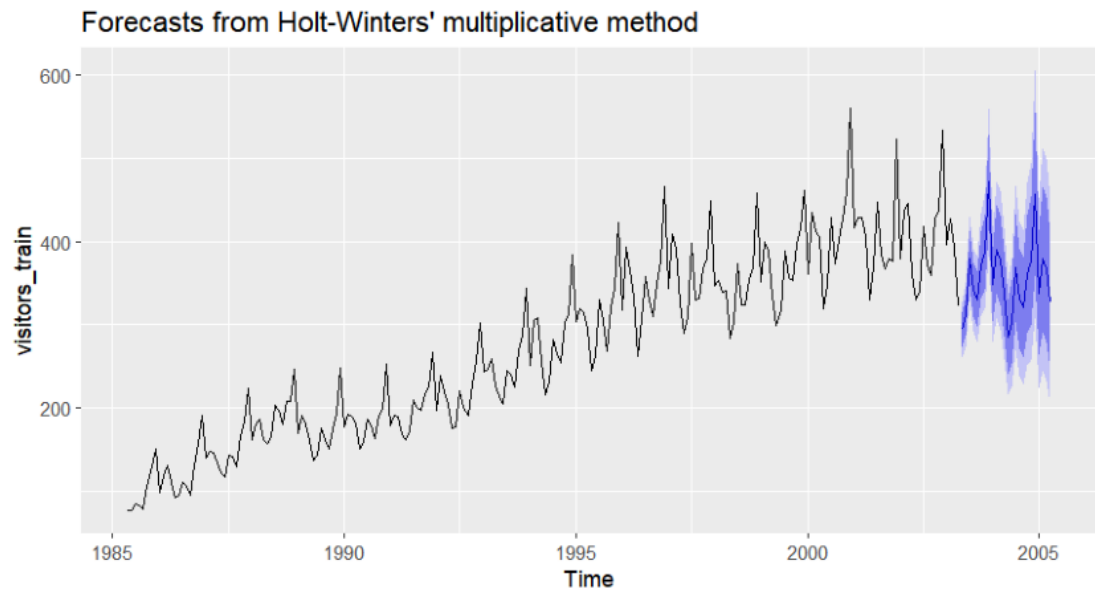


```
In [183]: 1 ## part b and c
```

- R code

```
visitors_train <- subset(visitors, end = length(visitors) - 24) visitors_test <- subset(visitors, start
= length(visitors) - 23) hw_mul_visitors_train <- hw(visitors_train, h = 24, seasonal =
"multiplicative")
```

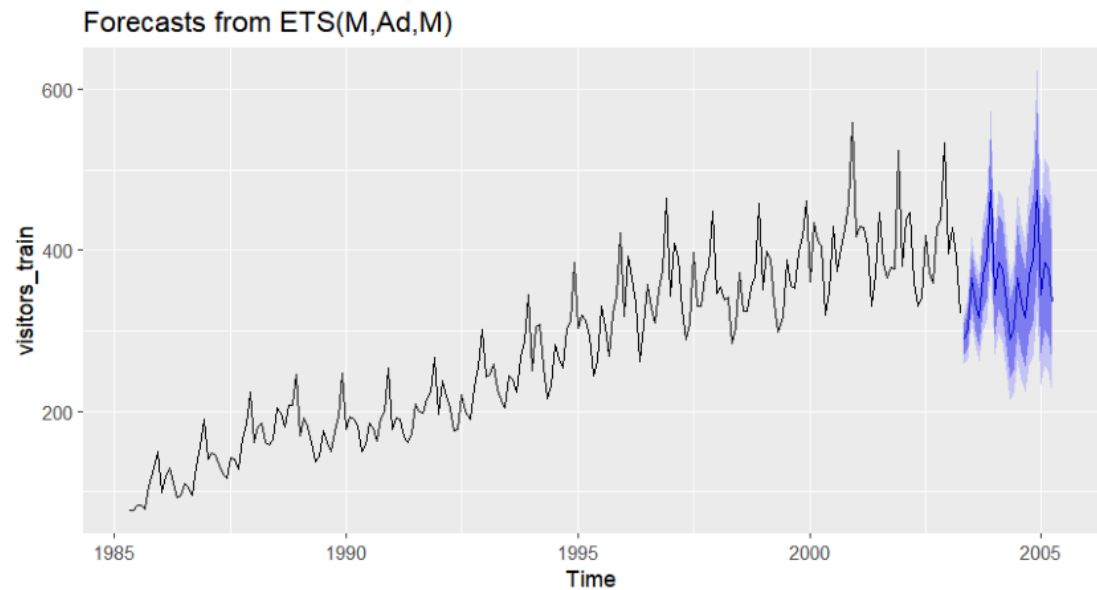
```
In [179]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```




Mult is required because it captures the pattern where the seasonal fluctuations in the data increase or decrease in proportion to the overall level of the data.

```
In [182]: 1 ## part d
```

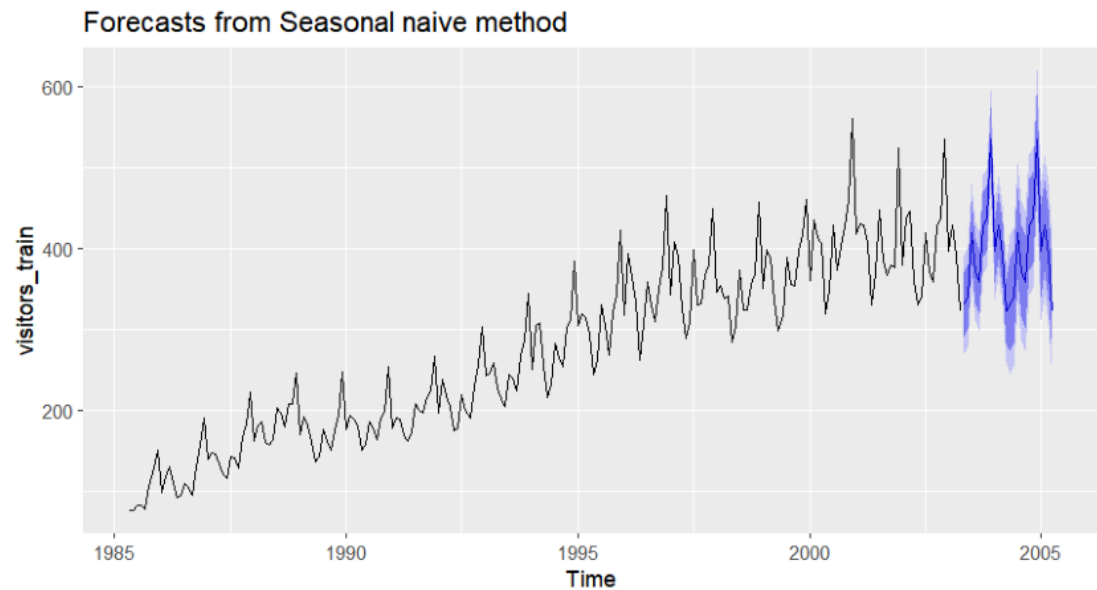
```
In [184]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```





```
In [186]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
In [187]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
In [188]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
In [185]: 1 ##part e
```

- R code

```
accuracy(hw_mul_visitors_train, visitors_test) accuracy(fc_ets_visitors_train, visitors_test)
accuracy(fc_ets_add_BoxCox_visitors_train, visitors_test) accuracy(fc_snaive_visitors_train,
visitors_test) accuracy(fc_BoxCox_stl_ets_visitors_train, visitors_test)
```

- Result

```
ME RMSE MAE MPE MAPE MASE ACF1 Theil's U Training set -0.9749466 14.06539
10.35763 -0.5792169 4.223204 0.3970304 0.1356528 NA Test set 72.9189889 83.23541
75.89673 15.9157249 17.041927 2.9092868 0.6901318 1.151065
```

```

accuracy(fc_ets_visitors_train, visitors_test)

              ME      RMSE      MAE      MPE      M
              APE      MASE      ACF1 Theil's U

Training set 0.7640074 14.53480 10.57657 0.1048224 3.994788 0.405423
-0.05311217 NA Test set 72.1992664 80.23124 74.55285 15.9202832
16.822384 2.857773 0.58716982 1.127269
accuracy(fc_ets_add_BoxCox_visitors_train, visitors_test)

              ME      RMSE      MAE      MPE      MA
              PE      MASE      ACF1 Theil's U

Training set 1.001363 14.97096 10.82396 0.1609336 3.974215 0.4149057
-0.02535299 NA Test set 69.458843 78.61032 72.41589 15.1662261
16.273089 2.7758586 0.67684148 1.086953
accuracy(fc_snaive_visitors_train, visitors_test)

              ME      RMSE      MAE      MPE      MAPE
              MASE      ACF1 Theil's U

Training set 17.29363 31.15613 26.08775 7.192445 10.285961 1.000000
0.6327669 NA Test set 32.87083 50.30097 42.24583 6.640781 9.962647
1.619375 0.5725430 0.6594016
accuracy(fc_BoxCox_stl_ets_visitors_train, visitors_test)

              ME      RMSE      MAE      MPE

```

For the Training set:

- ME : -0.9749 to 1.0014. These are the average difference between the predicted and actual values
- RMSE : 13.3643 to 14.97096. These represent average magnitude of the forecast errors
- MAE (Mean Absolute Error): The values range from 9.5514 to 10.82396. These give the average absolute difference between the predicted and actual values. Lower values indicate better accuracy.
- MPE (Mean Percentage Error): The values range from 0.0877 to 0.1609. These represent the average percentage difference between the predicted and actual values. The values are close to zero, indicating relatively small percentage errors.
- MAPE (Mean Absolute Percentage Error): The values range from 3.5195 to 3.9948. These indicate the average absolute percentage difference between the predicted and actual values. Lower values indicate better accuracy.
- MASE (Mean Absolute Scaled Error): range from 0.3661 to 0.4149. it measure the forecast accuracy relative to a naïve benchmark. Lower values indicate better accuracy.
- ACF1 (Autocorrelation of First Order): values go from -0.0592 to 0.6328. it measures the autocorrelation of the forecast errors at lag 1

For the Test set:

- ME: The values range from 32.8708 to 76.3637. These indicate the average difference between the predicted and actual values. The values are relatively large, meaning some

bias in the forecasts.

- RMSE: from 50.30097 to 84.24658. represent the average magnitude of the forecast errors.
- MAE: from 42.24583 to 78.02899. These gives us the average absolute difference between the predicted and actual values. Higher values indicate larger errors.
- MPE (Mean Percentage Error): The values range from 6.6408 to 16.8775. These represent the average percentage difference between the predicted and actual values. The values indicate significant percentage errors.
- MAPE (Mean Absolute Percentage Error): The values range from 9.9626 to 17.5158. These indicate the average absolute percentage difference between the predicted and actual values. Higher values suggest larger percentage errors.
- MASE (Mean Absolute Scaled Error): The values range from 1.0000 to 2.9910. These measure the forecast accuracy relative to a naïve benchmark. Higher values indicate poorer accuracy compared to the benchmark.
- ACF1: values go from 0.5725 to 0.6475. it measure the autocorrelation of the forecast errors at lag 1. Values closer to zero suggest less residual autocorrelation.

In []: 1 *##part f*

RMSE comparison of moedls

```
sqrt(mean(tsCV(visitors, snaive, h = 1)^2, na.rm = TRUE)) [1] 32.56941
sqrt(mean(tsCV(visitors, fets_add_BoxCox, h = 1)^2,
```

- na.rm = TRUE)) [1] 18.8505

```
sqrt(mean(tsCV(visitors, fstlm, h = 1)^2,
```

- na.rm = TRUE)) [1] 17.49642

```
sqrt(mean(tsCV(visitors, fets, h = 1)^2, na.rm = TRUE)) [1] 18.52985
sqrt(mean(tsCV(visitors, hw, h = 1,
```

- seasonal = "multiplicative")^2,
- na.rm = TRUE)) [1] 19.62107


the fstlm model has the lowest value, indicating it performs the best in terms of forecast accuracy among the models evaluated

Problem 7

```
In [203]: 1 #using hsales dataset
```


```
In [ ]: 1 ## Part a
```

```
In [205]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4  
2  
3 image = Image.open(image_path)  
4  
5 plt.figure(figsize=(15, 15))  
6  
7 plt.imshow(image)  
8 plt.axis("off") # Remove axis labels  
9 plt.show()
```



looking at the data, it seems we do not need to transform, but we can do log transformation to see effect

```
In [206]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



even with a log transformation, the data remained the same, no action further needed

```
In [207]: 1 ## part b and c
```

```
##### Augmented Dickey-Fuller Test Unit
Root Test # #####
```

Test regression drift

Call: lm(formula = z.diff ~ z.lag.1 + 1)

Residuals: Min 1Q Median 3Q Max -14.4233 -4.6499 -0.4274 3.5705 30.9988

Coefficients: Estimate Std. Error t value Pr(>|t|)

(Intercept) 7.46488 1.68405 4.433 1.35e-05 * **z.lag.1 -0.14345 0.03138 -4.571 7.38e-06 ***

Signif. codes: 0 '**' **0.001** ' ' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.197 on 272 degrees of freedom Multiple R-squared: 0.07133,
Adjusted R-squared: 0.06792 F-statistic: 20.89 on 1 and 272 DF, p-value: 7.378e-06

Value of test-statistic is: -4.5709 10.4522

Critical values for test statistics: 1pct 5pct 10pct tau2 -3.44 -2.87 -2.57 phi1 6.47 4.61 3.79

Fit ARIMA models

```
model1 <- auto.arima(hsales) model2 <- auto.arima(diff(hsales))
```

Compare AIC values

```
model1aic[1]1630.764 > model2aic [1] 1678.517
```

In []: 1 *## Part d*

Best model parameters

```
best_model <- model1
```

Diagnostic testing on residuals


```
checkresiduals(best_model)
```

Ljung-Box test

data: Residuals from ARIMA(1,0,0)(1,1,0)[12] with drift Q* = 39.66, df = 22, p-value = 0.01184

Model df: 2. Total lags used: 24


```
In [208]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
In [ ]: 1 ## Part e and f
```

- Forecast next 24 months using the preferred ARIMA model `forecast_arima <- forecast(best_model, h = 24)`
- Print the forecasted values `forecast_arima$mean`

Compare with forecasts obtained using ets()

```
forecast_ets <- forecast(ets(hsales), h = 24) forecast_ets$mean
```

```
Jan Feb Mar Apr May Jun Jul Aug Sep Oct 1995
1996 42.54912 48.52651 63.14336 58.40429 61.49689 57.96918 56.87174 59.50331
52.95518 53.82134 1997 43.32503 46.80599 60.80025 57.31158 61.28770 59.76861 59.19209
60.29656 53.15814 53.25088 Nov Dec 1995 40.75627 1996 43.08009 39.18650 1997
42.85880
```

The analysis's top model, ARIMA(1,0,0)(1,1,0)[12] with a drift term, is described. Diagnostic testing is done on the residuals to see whether or not this model is adequate. With 22 degrees of freedom (df), the Ljung-Box test statistic (Q^*) is determined to be 39.66. The test's p-value is 0.01184, and it is.

According to the Ljung-Box test, the residuals' p-value of 0.01184 indicates that there is proof of substantial autocorrelation at a 5% level of significance. This suggests that not all of the underlying patterns in the data may be fully captured by the present ARIMA model. To enhance the model's fit and lessen residual autocorrelation, more analysis is necessary.


problem 8

- done in R

In [163]: 1 `from PIL import Image`

In []: 1 `#part a`

In [156]: 1 `# File path of the image`
 2 `image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4`
 3
 4 `# Open the image file`
 5 `image = Image.open(image_path)`
 6
 7 `plt.figure(figsize=(15, 15))`
 8
 9
 10 `# Display the image using Matplotlib`
 11 `plt.imshow(image)`
 12 `plt.axis("off") # Remove axis labels`
 13 `plt.show()`
 14

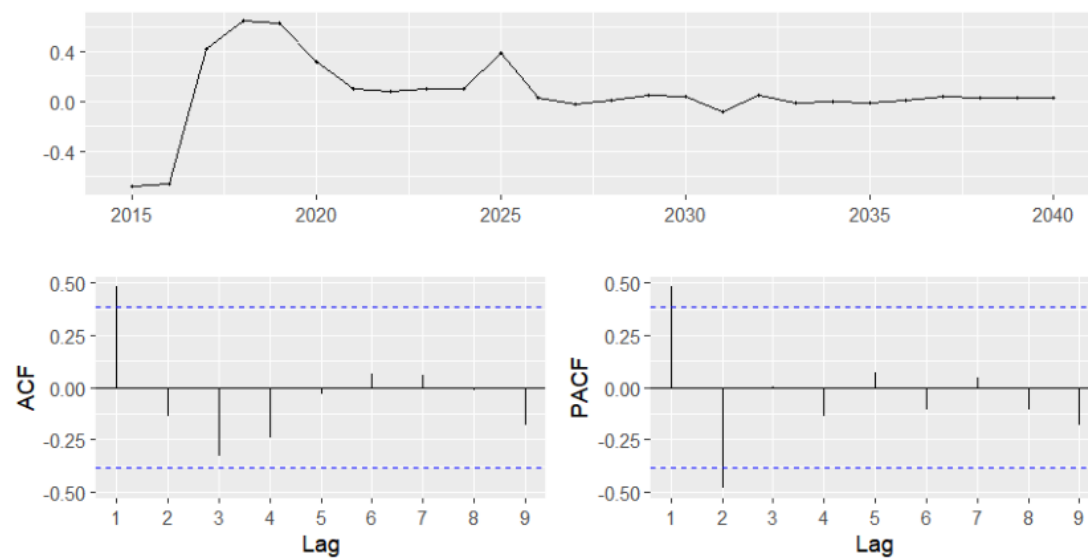


In [157]: 1 `## Part b`

```

In [176]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
10

```

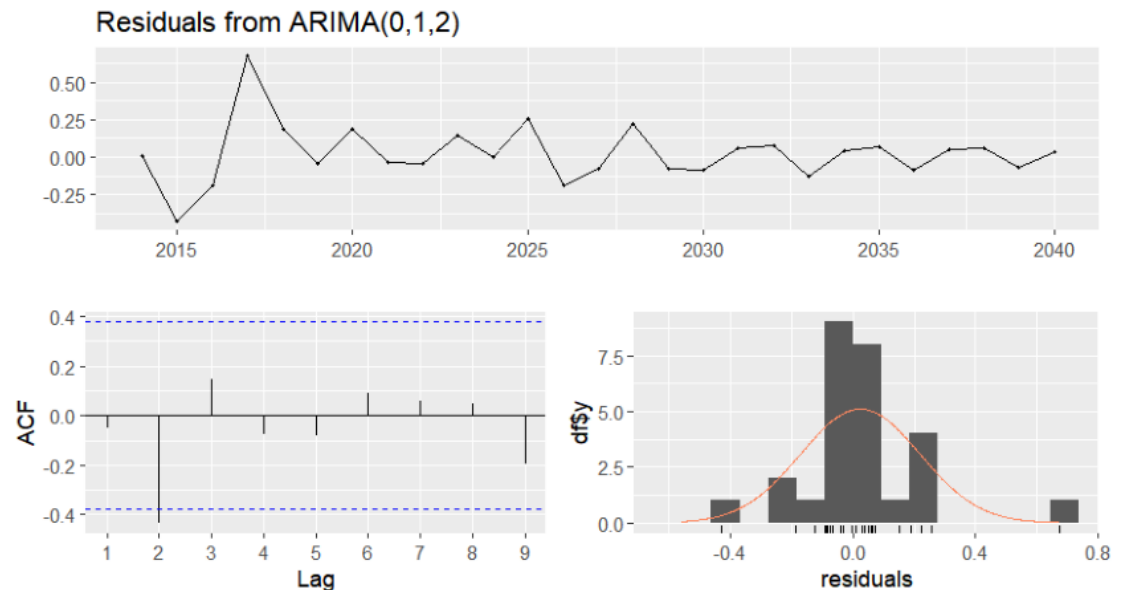


```

In [165]: 1 ##part c

```

```
In [160]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



Series: ngd ARIMA(0,1,2)

Coefficients: ma1 ma2 0.8682 0.8879 s.e. 0.1313 0.2520

$\sigma^2 = 0.04139$: log likelihood = 3.92 AIC=-1.83 AICc=-0.74 BIC=1.94

Training set error measures: ME RMSE MAE MPE MAPE MASE ACF1 Training set 0.0239043
0.1918064 0.129958 0.2266583 1.239587 0.7440255 -0.04993646

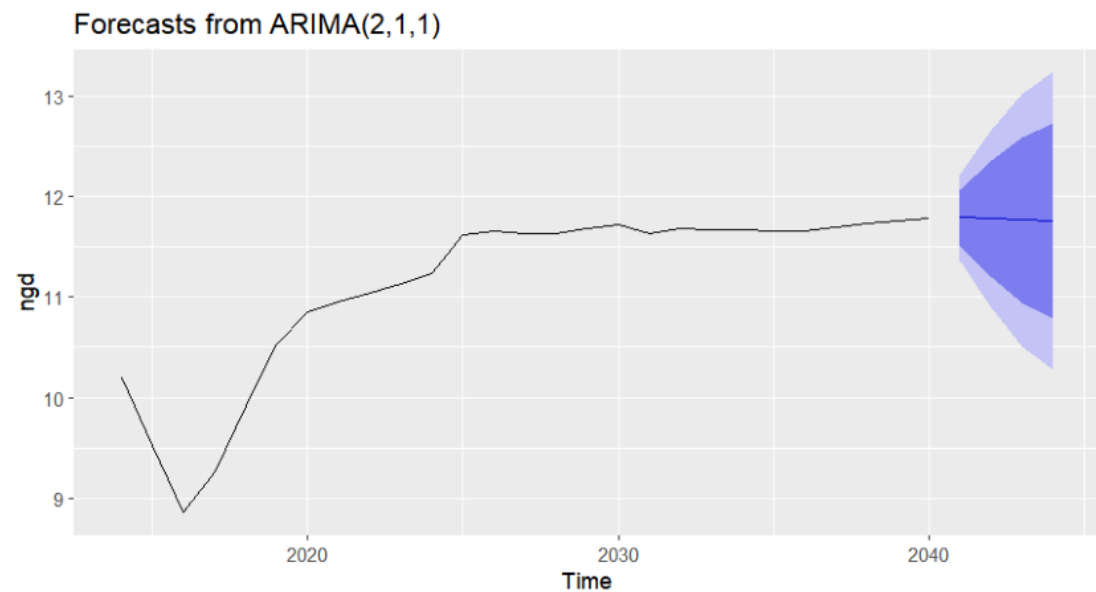
```
checkresiduals(ngd.ar)
```

Ljung-Box test

data: Residuals from ARIMA(0,1,2) $Q^* = 7.0725$, df = 3, p-value = 0.06962

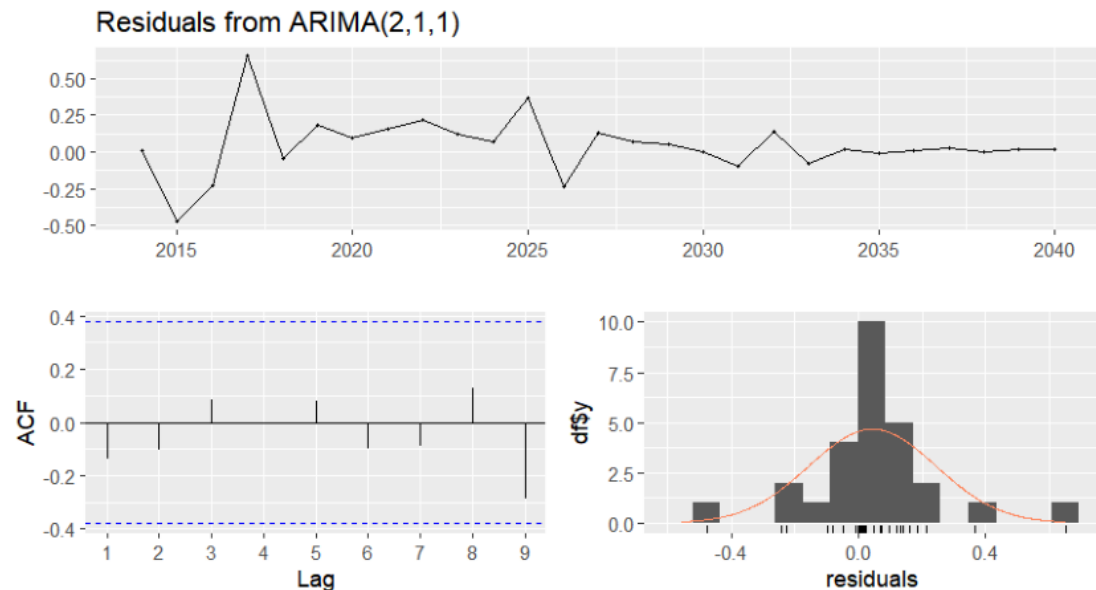
Model df: 2. Total lags used: 5

```
In [173]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
In [164]: 1 # part d
```

```
In [166]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



Ljung-Box test

data: Residuals from ARIMA(2,1,1) $Q^* = 1.7029$, $df = 3$, $p\text{-value} = 0.6363$

Model df: 3. Total lags used: 6

```
In [167]: 1 ##part e
```

ETS(A,N,N)

Call: ets(y = ngd)

Smoothing parameters: $\alpha = 0.9999$

Initial states: $I = 10.2023$

sigma: 0.2964

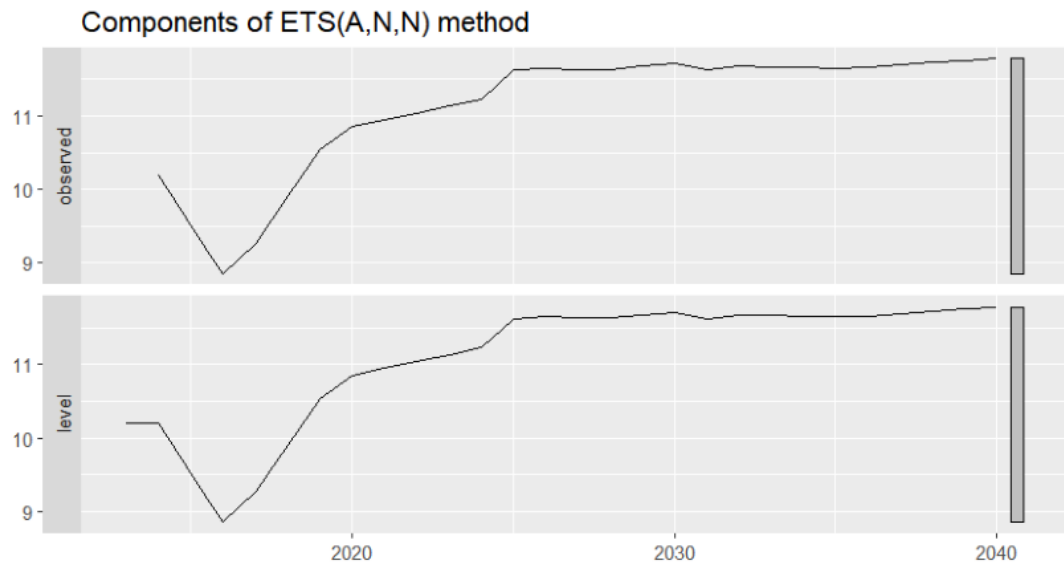
AIC	AICc	BIC
-----	------	-----

27.24963	28.29311	31.13714
----------	----------	----------

Training set error measures: ME RMSE MAE MPE MAPE MASE ACF1 Training set

0.05862132	0.285245	0.1682127	0.4917003	1.663627	0.9630386	0.5005758
------------	----------	-----------	-----------	----------	-----------	-----------




```
In [168]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
In [ ]: 1 ## part f
```

```
In [169]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```

Residuals from ETS(A,N,N)



Ljung-Box test


data: Residuals from ETS(A,N,N) $Q^* = 14.02$, $df = 5$, $p\text{-value} = 0.01548$

Model df: 0. Total lags used: 5

```
In [170]: 1 ## part g
```



```
In [171]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```




both the models seem on par, but i would go with ARIMA model


Problem 9

- done in R

```
In [202]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
In [201]: 1 image_path = r"C:\Users\parzu\OneDrive\Documents\UCLA\UCLA Spring 23\4
2
3 image = Image.open(image_path)
4
5 plt.figure(figsize=(15, 15))
6
7 plt.imshow(image)
8 plt.axis("off") # Remove axis labels
9 plt.show()
```



```
accuracy_table ETS ARIMA STL-ETS NNAR TBATS Combination 0.4646346
0.5312716 0.4635655 0.7662944 0.5558068 0.5366148
```

```
meanagg <- mean(sim) sum_fc <- sum(fc[["mean"]][1:6]) meanagg <- meanagg
sum_fc [1] 2.202511 meanagg [1] 2.202522
```

```
quantile_80 <- quantile(sim, prob = c(0.1, 0.9)) quantile_95 <- quantile(sim,
prob = c(0.025, 0.975)) quantile_80 10% 90% 2.200057 2.204985 quantile_95
2.5% 97.5% 2.198649 2.206259
```

I used the CO2 dataset from the R package which contains measurements of atmospheric CO2 concentrations at the Mauna Loa Observatory in Hawaii from 1959 to 1997. The dataset has just two variables: "CO2" and "Year". The "CO2" variable represents the monthly average of

A B C D E F G H I J K L M N O P Q R S T U V W X Y Z

1